



# Programming Assignment II

## COSI 127B: Introduction to Database Systems

Eden Zik

Brandeis University

2015

# Introduction

- ▶ PA1 covered skills needed to *query and modify* a database
  - ▶ Connecting to Postgres and initializing tables.
  - ▶ Inserting and deleting data.
  - ▶ Composing complex queries
- ▶ PA2 covers skills needed to *manage and integrate* a database
  - ▶ Writing integrity constraints
  - ▶ Implementing triggers
  - ▶ Using SQL with Java to power *data driven applications*

First of all, some background...

# Background

- ▶ The TPC-H Benchmark
- ▶ Integrity Constraints
- ▶ Triggers
- ▶ JDBC

# The TPC-H Benchmark

- ▶ A decision support benchmark for database systems.
  - ▶ Consists of a schema, data, and a standard suite of queries
  - ▶ Models the data needs of a manufacturing company
  - ▶ Used to compare performance of different database systems
- ▶ Schema
  - ▶ Eight tables populated with data.
  - ▶ Each attribute is prefixed - unique across the database

<b>Table Name</b>	<b>Prefix</b>	<b>Primary Key</b>
region	r	regionkey
nation	n	nationkey
supplier	s	suppkey
part	p	partkey
partsupp	ps	partkey, suppkey
customer	c	custkey
lineitem	l	orderkey, linenumbr
orders	o	orderkey

Next up: Integrity Constraints

# Integrity Constraints<sup>1</sup>

- ▶ Four types of constraints
  - ▶ Key Constraints
  - ▶ Attribute Constraints
  - ▶ Referential Integrity Constraints
  - ▶ Global Constraints
- ▶ Syntax
  - ▶ `ALTER TABLE mytable ADD PRIMARY KEY (thekey)`
  - ▶ `ALTER TABLE mytable  
ADD CONSTRAINT price CHECK (price > 0)`
  - ▶ `ALTER TABLE mytable  
ADD CONSTRAINT price NOT NULL`
  - ▶ `ALTER TABLE mytable  
ADD CONSTRAINT price UNIQUE`
  - ▶ `ALTER TABLE mytable  
ADD CONSTRAINT fk1 FOREIGN KEY (otherkey)  
references yourtable`

When Integrity Constraints aren't enough: Triggers

---

<sup>1</sup>Covered in Lecture 6

# Triggers

- ▶ Callback functions defined over tables or views
- ▶ Executed whenever a certain type of operation is performed
  - ▶ INSERT
  - ▶ UPDATE
  - ▶ DELETE
- ▶ Can be used to enforce sophisticated integrity constraints
  - ▶ Propagating value change
    - ▶ Ex. "When sales tax rate changes, update all prices"
    - ▶ Why not use an UPDATE in a loop? Establishes an internal dependency between prices and sales tax that is easier to maintain
  - ▶ Rejecting modifications leading to illegal states not otherwise enforceable by simpler constraints.
    - ▶ Ex. "A family plan cannot have more than 6 members"
    - ▶ Why not use global constraints? Too expensive

# Triggers cont.

- ▶ Defined using PL/pgSQL
  - ▶ PostgreSQL specific procedural language
  - ▶ SQL with some programming facilities
    - ▶ Accepts arguments
    - ▶ Executes arbitrary SQL
    - ▶ Can return a typed value (in this case, TRIGGER).
  - ▶ 

```
CREATE FUNCTION some_function()  
RETURNS TRIGGER AS $$  
BEGIN  
INSERT INTO some_table(some_att) VALUES ('value');  
RETURN NEW;  
END  
$$ LANGUAGE 'plpgsql';
```

## Triggers cont.

- ▶ Added on a table by CREATE TRIGGER syntax
  - ▶ Defined BEFORE, AFTER, or INSTEAD OF an event
  - ▶ Events can be INSERT, UPDATE, or DELETE's
  - ▶ Executes PROCEDURE when event occurs either FOR EACH ROW or FOR EACH STATEMENT
  - ▶ 

```
CREATE TRIGGER some_trigger  
AFTER INSERT ON some_table  
FOR EACH ROW EXECUTE PROCEDURE some_function();
```

Finally: JDBC



# JDBC

- ▶ Previously...
  - ▶ Issued queries via a Command Line Interface
  - ▶ Viewed query results as texts
- ▶ In practice...
  - ▶ Queries issued programmatically
    - ▶ JDBC = [J]ava [D]ata[b]ase [C]onnectivity
    - ▶ Enables connecting to a database via Java
    - ▶ Manipulate query results (ResultSet) like native data structures
  - ▶ Results of queries power applications
    - ▶ Interactivity
    - ▶ Data visualization
    - ▶ Report generation
  - ▶ We will use JDBC to visualize trends in the TPC-H database and generate reports based on part orders.

# Example of a JDBC Program

```
import java.sql.*;
public class JDBCExample {
    static final String DB_TYPE = "postgresql";
    static final String DB_DRIVER = "jdbc";
    static final String DB_NAME = System.getenv("PGDATABASE");
    static final String DB_HOST = System.getenv("PGHOST");
    static final String DB_URL = String.format("%s:%s://%s/%s", DB_DRIVER, DB_TYPE,
                                                DB_HOST, DB_NAME);

    static final String DB_USER = System.getenv("PGUSER");
    static final String DB_PASSWORD = System.getenv("PGPASSWORD");
    static Connection conn;

    static final String QUERY = "SELECT_'hello_world!';";

    public static void main(String[] args) throws SQLException {
        conn = DriverManager.getConnection(DB_URL, DB_USER, DB_PASSWORD);
        Statement st = conn.createStatement();
        ResultSet rs = st.executeQuery(QUERY);
        rs.next();
        System.out.println(rs.getString(1));
    }
}
```

# Using JDBC

- ▶ Set database connection parameters
- ▶ Initialize a `Connection` object
- ▶ Initialize a `Statement` object using the connection
- ▶ Fill a `ResultSet` object by calling `executeQuery` on a `Statement`
- ▶ Iterate through each row in `ResultSet`
- ▶ Fetch column by number or name from the current row

So what will we do in PA2?

# Overview

- Part 1
  - ▶ Use triggers and the integrity constraints to enforce specific semantics on the TPC-H database.
  - ▶ Propagate changes in the data via the use of triggers.
- Part 2
  - ▶ Produce visualizations based on TPC-H data using JDBC and `gnuplot/graphviz`
  - ▶ Generate purchase reports based on TPC-H part orders using JDBC and  $\text{\LaTeX}$ .

# Getting Started: Environment

- ▶ Log on to a CS Public Macine using SSH

```
$ ssh [your username]@[cs public machine name].cs.brandies.edu
```

- ▶ Copy the provided code into your home directory

```
$ cp /home/o/class/cs127b/PA2/PA2Files.tar.gz .
```

```
$ tar -xvzf PA2Files.tar.gz
```

# Getting Started: Database

- ▶ Set up your environment variables to connect

```
$ export PGHOST=[your database host]
```

```
$ export PGUSER=[your user name]
```

```
$ export PGDATABASE=[your user name]pa2
```

```
$ export PGPASSWORD=[your database password]
```

- ▶ Initialize the database

```
$ make initialize-db
```

- ▶ Use `psql` to test all tables have been loaded

- ▶ Reset the database as necessary

```
$ make reset-db
```

# Getting Started: JDBC

- ▶ Familiarize yourself with the provided files
  - ▶ `makefile` - A makefile to ease setup and compilation
  - ▶ `README` - this list of files
  - ▶ `lib/jdbc.jar` - The Postgres database driver
  - ▶ `lib/pa2.jar` - API's for Gnuplot, Graphviz, and LaTeX
  - ▶ `docs/index.html` - Javadoc documentation for said API's
  - ▶ `src/part1.java` - Template code for Part 2.1
  - ▶ `src/part2.java` - Template code for Part 2.2
  - ▶ `src/part3.java` - Template code for Part 2.3
  - ▶ `txt/orders.txt` - Purchase orders for part 2.3
- ▶ Make sure your database is initialized and your environment variables set


## Getting Started: Part 1.\*

- ▶ Part 1.1 will require you to write primary and foreign key constraints. All of these should be saved as files under the following directory structure:  
`./sql/keys/[primary / foreign].sql`
- ▶ Part 1.2 will require you to write constraints and triggers. All of these should be saved as files under the following directory structure:  
`./sql/triggers/q-[question number].sql`




# Getting Started: Part 2.1

- ▶ Part 2.1 will require you to use an API to Gnuplot
  - ▶ Gnuplot → plotting library
  - ▶ `import edu.brandeis.cs127b.pa2.gnuplot.*;`
  - ▶ Documented in `docs/*`
  - ▶ Starter code produces a simple plot ( $\sin(x)$ )
  - ▶ Your task: *calculate supplier sales totals by region*



## Getting Started: Part 2.2

- ▶ Part 2.2 will require you to use an API to Graphviz
  - ▶ Graphviz → graph library
  - ▶ `import edu.brandeis.cs127b.pa2.graphviz.*;`
  - ▶ Documented in `docs/*`
  - ▶ Starter code produces a simple graph (random edges)
  - ▶ Your task: *calculate the total sales between suppliers and customers grouped by the regions where the suppliers and customers are located*



# Getting Started: Part 2.3

- ▶ Part 2.3 will require you to use an API to L<sup>A</sup>T<sub>E</sub>X
  - ▶ L<sup>A</sup>T<sub>E</sub>X → document typesetting library
  - ▶ `import edu.brandeis.cs127b.pa2.latex.*;`
  - ▶ Documented in `docs/*`
  - ▶ Starter code produces a simple document (empty order)
  - ▶ Your task: *determines minimum prices for each of a specified set of items in a file, and generate a purchase order that orders each item from its cheapest supplier*

## Brandeis University

415 South St. (781) 736-2000  
Waltham, MA 02453

Date:

10/24/2015

Order #:

4

Part #	Quantity	Unit Price	Amount
Supplier: 106			
1223	95	\$253.55	\$24087.25
Subtotal			\$24087.25
Supplier: 108			
8911	37	\$344.32	\$12729.84
Subtotal			\$12729.84
Supplier: 109			
1822	50	\$53.60	\$2680.00
3095	12	\$116.39	\$1396.68
3229	59	\$302.26	\$23133.34
8714	4	\$14.38	\$57.52
Subtotal			\$27277.54
Supplier: 11			
1206	15	\$82.20	\$1233.00
679	4	\$177.54	\$710.16
Subtotal			\$1943.16
Total			\$66047.79
# of Suppliers			4

## Getting Started: Part 2.\*

- ▶ Compile your code in `src/*` as follows:  
`$make build-[part1 / part2 / part3] → ./bin`
- ▶ Run your code as follows:  
`$make run-[part1 / part2 / part3] →  
./[plot/gv/tex]`
- ▶ Visualize your result as follows:  
`$make vis-[part1 / part2 / part3] → ./vis`
- ▶ Submit a well commented source code in `./src`, along with your `gnuplot`, `graphviz`, and `LATEX` code and their resulting visualizations
- ▶ A detailed writeup about your algorithm for each part in a PDF format in `./pdf`.

# Submission

1. Set the proper environment variables:

```
$ export FIRSTNAME=[your first name as it appears on LATTE]
```

```
$ export LASTNAME=[your last name as it appears on LATTE]
```

2. Type the following command to package all relevant files into a submittable file:

```
$ make submit
```

A file with the name [your last name]\_[your first name]\_127b\_pa2.tar.gz will be created in the current directory.

3. Use secure copy to transfer the file from the server to your machine:

```
$ scp [your username]@[cs public machine  
name].cs.brandies.edu: /[pa2]/*_*_127b_pa2.tar.gz .
```

The file should then be transferred to the current directory on your local machine

4. Upload the file to LATTE.